

## CS 380D - DISTRIBUTED COMPUTING RAFT - PROJECT REPORT

### PROJECT STRUCTURE:

#### **config.py**

- Manages the system configurations
- Contains NodeState enumeration which defines the states that a node can be in - FOLLOWER, CANDIDATE, and LEADER
- Sets the base port and host for server communication
- Defines the timing constraints for leader elections and heartbeat
- Manages the membership of nodes in the cluster; stores the list and the number of active nodes in the cluster in JSON files
  - save\_servers():
    - save the IDs, ports, and addresses of the servers in the system
  - save\_cluster\_members():
    - saves a list of active servers in the system

#### **raft\_server.py**

- Initializes a RaftNode instance with the given node\_id as **raftserver{servername}**
- Starts the RAFT node, which:
  - initializes the gRPC server
  - calls config to add the server to the list
- Keeps the process running until interrupted
- Upon receiving a shutdown signal, cleans up resources, stops the gRPC server, and calls config to update the cluster configuration

#### **state\_manager.py**

- StateManager class initializes the node's state by either loading it from a file or creating a new state structure
- State structure:
  - current\_term: server's term
  - voted\_for: candidate for which the server voted
  - log:

- term in which entry was created
- log index
- command (PUT, GET, REPLACE)
- key, and
- value to be stored in the key-value store
- commit\_index: highest log index that was committed
- last\_applied: highest log index applied to the state machine
- **data: key-value store**
- sent\_length: tracks the last log index sent to each follower
- acked\_length: tracks the highest log index acknowledged by each follower
- request\_history: tracks processed client requests to avoid duplicates
- StateManager class also contains getters and setters for the state variables
- is\_duplicate\_request():
  - serves as an idempotent function to check for duplicate requests
- record\_request():
  - stores the result of processed request
- append\_log\_entry():
  - creates an entry to be appended
  - checks for conflicting entries in server log and deletes them
- commit\_logs\_up\_to():
  - commits the key-value pair into the data store from last applied to index
  - commits the client and request id pairs for duplicity check

### **raft.proto**

- Specifies RPC methods for inter-server and client communication
- Defines messages that are passed as requests and responses during communication
- The methods were defined based on the specifications mentioned

### **node.py**

- ConnectionManager class manages gRPC connections for internal server communication (7000 ports)
  - \_calculate\_source\_port():

- maintains unique source ports for each server
  - `get_channel()`:
    - reuses and creates channels from source ports (7000s) to target ports (9000s)
  - `close_all()`:
    - closes all connections
- `KeyValueStoreService` class acts as the intermediary between frontend and raft node
  - `GET/PUT/REPLACE`:
    - converts requests to suit the input format of the `SendCommand` function
  - `_convert_response()`:
    - converts responses from `SendCommand` to suit frontend response format
  - `GetState()`:
    - returns the state (leader) and term of the server
- `RaftNode` class holds the attributes of a server like `node_id`, status, logs, terms, and data
- Manages state transitions and leader elections
  - `_run_election_timer()`:
    - sends heartbeats
    - triggers leader elections if heartbeats have not been received
  - `_start_election()`:
    - when timeout occurs, changes the state of the server to candidate and requests votes from the other servers by sending `RequestVote()` RPC
- Manages log replication
  - Sends `AppendEntries` RPC to replicate log so that each server has an up-to-date log
  - `_send_heartbeats()`:
    - sends empty `AppendEntries()` to inform that the leader is active and also to inform the commit index
- Handles client requests and ensures consistency
- Multiple other functions to ensure the running and proper functioning of the raft nodes

## **frontend.py**

- FrontEndService as an interface for communication with clients
- Manages RAFT lifecycle - starting and stopping servers, and cleaning up
  - StartRaft():
    - starts a cluster with a specific number of servers
    - assigns sequential ports starting from the base port
  - \_cleanup\_servers():
    - terminates raft servers processes and ensures there is no garbage
- \_find\_leader():
  - determines the current leader by querying other servers by using gRPC
  - implements gRPC methods which enable the client to communicate (GET, PUT, REPLACE) with the servers
  - defines a forwarding mechanism to enable the requests to be passed to the leader in case the clients contact the follower nodes

## **IMPLEMENTATION DETAILS:**

### **● Port Assignments**

Each node calculates its source port using the formula:

Source Port = 7000 + (Server ID - 1) \* Cluster Size + Target Server Offset

Target server offset is adjusted to skip the node's own ID

### **● Process Naming**

Each RAFT server process is named as raftserver<ID> for easy identification:

```
setproctitle.setproctitle(f"raftserver{server_num}")
```

### **● State Persistence**

State is saved to JSON files for durability:

with open(self.state\_file, 'w') as f:

```
    json.dump(self.state, f)
```

- **Log Replication**

Entries are appended and replicated to followers along with the heartbeat signal

- **Leader Election**

Nodes initiate elections on timeout (when they don't receive a heartbeat)  
if `time_since_last_heartbeat > self.election_timeout`:  
    `self._start_election()`

Votes are requested and tallied based on acknowledgments from a quorum:  
`response = stub.RequestVote(request)`  
if `response.vote_granted`:  
    `self.votes_received += 1`

- **Client Request Handling**

Requests are forwarded to the leader if the clients contact follower nodes:  
`stub = self._get_stub_for_server(self.leader_id)`  
`response = stub.Put(request)`

## **DEPENDENCIES:**

Mentioned in requirements.txt

## **NOTE:**

- We might have to run frontend.py a few times to get the right output due to the randomness in timeouts for different nodes.
- **Since Python's gRPC does not bind a server to a port the way C, C++, Go, etc., do, blocking a port using an IP table does not ensure that the inter-node communication is blocked, since the server might not be using the port. So, we attempted to simulate blocking of nodes by killing them, which is in mytest.go file.**
- Additionally, we have added a timeout of 1 second after every PUT in the loadDataset() in mytest.go file.